

## Problem Statement

SONAR uses sound signals from submarine needs to detect that object underwater is rock or mine

## Work Flow

1. Sonar Data:- labortary data is used here which is obtained from rock and metal cylinder 2. Data Preprocessing:- need to do analyse the and make data fit for modelling 3. Split Data:- splitting of data into test-train 4. Model Selection:- choosing the best model 5. Model Evaluation:- check for the model evaluation

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [2]: df=pd.read_csv('sonar_data.csv',header=None)
```

```
In [3]: df
```

```
Out[3]:
```

	0	1	2	3	4	5	6	7	8	9	...	51	52	53	54	55	
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159	0.0072	0.0167	0.0
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048	0.0094	0.0191	0.0
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095	0.0180	0.0244	0.0
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150	0.0085	0.0073	0.0
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105	0.0110	0.0015	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
203	0.0187	0.0346	0.0168	0.0177	0.0393	0.1630	0.2028	0.1694	0.2328	0.2684	...	0.0116	0.0098	0.0199	0.0033	0.0101	0.0
204	0.0323	0.0101	0.0298	0.0564	0.0760	0.0958	0.0990	0.1018	0.1030	0.2154	...	0.0061	0.0093	0.0135	0.0063	0.0063	0.0
205	0.0522	0.0437	0.0180	0.0292	0.0351	0.1171	0.1257	0.1178	0.1258	0.2529	...	0.0160	0.0029	0.0051	0.0062	0.0089	0.0
206	0.0303	0.0353	0.0490	0.0608	0.0167	0.1354	0.1465	0.1123	0.1945	0.2354	...	0.0086	0.0046	0.0126	0.0036	0.0035	0.0
207	0.0260	0.0363	0.0136	0.0272	0.0214	0.0338	0.0655	0.1400	0.1843	0.2354	...	0.0146	0.0129	0.0047	0.0039	0.0061	0.0

208 rows × 61 columns



```
In [4]: df.isnull().sum()
```

```
Out[4]: 0      0
1      0
2      0
3      0
4      0
..
56     0
57     0
58     0
59     0
60     0
Length: 61, dtype: int64
```

observation:- no null values in the data

```
In [5]: df.duplicated().sum()
```

```
Out[5]: 0
```

observation:- no duplicate value in the dataset

```
In [6]: df[60].value_counts()
```

```
Out[6]: 60
M      111
R       97
Name: count, dtype: int64
```

no need to perform under sampling

observation:- signal has detected 111 mines and 97 rocks in labortary experiment

```
In [7]: #checking if their is need for standardization or not
df.drop(60,axis=1).std()
```

```
Out[7]: 0      0.022991
        1      0.032960
        2      0.038428
        3      0.046528
        4      0.055552
        5      0.059105
        6      0.061788
        7      0.085152
        8      0.118387
        9      0.134416
       10      0.132705
       11      0.140072
       12      0.140962
       13      0.164474
       14      0.205427
       15      0.232650
       16      0.263677
       17      0.261529
       18      0.257988
       19      0.262653
       20      0.257818
       21      0.255883
       22      0.250175
       23      0.239116
       24      0.244926
       25      0.237228
       26      0.245657
       27      0.237189
       28      0.240250
       29      0.220749
       30      0.213992
       31      0.213237
       32      0.206513
       33      0.231242
       34      0.259132
       35      0.264121
       36      0.239912
       37      0.212973
       38      0.199075
       39      0.178662
       40      0.171111
       41      0.168728
       42      0.138993
       43      0.133291
       44      0.151628
       45      0.133938
       46      0.086953
       47      0.062417
       48      0.035954
       49      0.013665
       50      0.012008
       51      0.009634
       52      0.007060
       53      0.007301
       54      0.007088
       55      0.005736
       56      0.005785
       57      0.006470
       58      0.006181
       59      0.005031
dtype: float64
```

no need to perform standartization as all the data is in common format and range(standard deviation is around 1)

```
In [8]: #statistical analysis
df.describe()
```

Out[8]:		0	1	2	3	4	5	6	7	8	9	..
	<b>count</b>	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	..
	<b>mean</b>	0.029164	0.038437	0.043832	0.053892	0.075202	0.104570	0.121747	0.134799	0.178003	0.208259	..
	<b>std</b>	0.022991	0.032960	0.038428	0.046528	0.055552	0.059105	0.061788	0.085152	0.118387	0.134416	..
	<b>min</b>	0.001500	0.000600	0.001500	0.005800	0.006700	0.010200	0.003300	0.005500	0.007500	0.011300	..
	<b>25%</b>	0.013350	0.016450	0.018950	0.024375	0.038050	0.067025	0.080900	0.080425	0.097025	0.111275	..
	<b>50%</b>	0.022800	0.030800	0.034300	0.044050	0.062500	0.092150	0.106950	0.112100	0.152250	0.182400	..
	<b>75%</b>	0.035550	0.047950	0.057950	0.064500	0.100275	0.134125	0.154000	0.169600	0.233425	0.268700	..
	<b>max</b>	0.137100	0.233900	0.305900	0.426400	0.401000	0.382300	0.372900	0.459000	0.682800	0.710600	..

8 rows × 60 columns

```
In [9]: df.groupby(60).mean()
```

Out[9]:		0	1	2	3	4	5	6	7	8	9	...	50	51
	<b>60</b>													
	<b>M</b>	0.034989	0.045544	0.050720	0.064768	0.086715	0.111864	0.128359	0.149832	0.213492	0.251022	...	0.019352	0.016014
	<b>R</b>	0.022498	0.030303	0.035951	0.041447	0.062028	0.096224	0.114180	0.117596	0.137392	0.159325	...	0.012311	0.010453

2 rows × 60 columns

```
In [10]: #separating features and target variable
x=df.drop(60,axis=1)
```

```
In [11]: y=df[60]
```

```
In [12]: x.shape
```

```
Out[12]: (208, 60)
```

```
In [13]: y.shape
```

```
Out[13]: (208,)
```

```
In [14]: #splitting the data into train and test
from sklearn.model_selection import train_test_split
```

```
In [15]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2,stratify=y)
#stratify-> data will be split acc. to no. of rock and mine
```

model used for evaluation 1.logistic 2.decision tree 3.random forest 4.naive bayes 5.knn

```
In [16]: #logistic
from sklearn.linear_model import LogisticRegression
```

```
In [17]: reg=LogisticRegression()
```

```
In [19]: reg.fit(x_train,y_train)
```

```
Out[19]: LogisticRegression
LogisticRegression()
```

```
In [20]: y_predict=reg.predict(x_test)
```

```
In [21]: #model evaluation
from sklearn.metrics import accuracy_score
```

```
In [24]: logistic_accuracy=round(accuracy_score(y_predict,y_test),3)
logistic_accuracy
```

```
Out[24]: 0.833
```

```
In [25]: #decision tree
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
```

```
In [26]: decision_model=DecisionTreeClassifier()
```

```
In [27]: decision_model.fit(x_train,y_train)
```

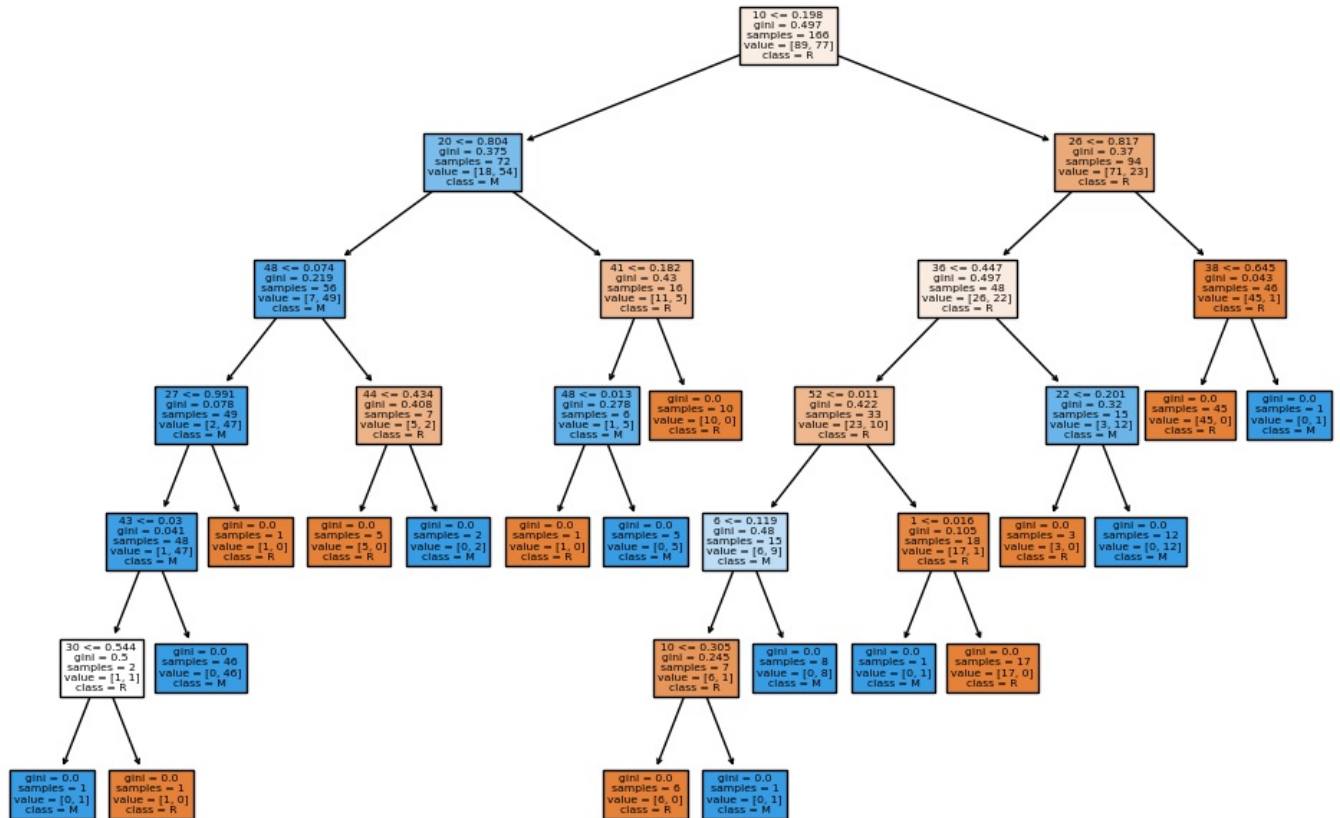
```
Out[27]: ▼ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [28]: y_predict=decision_model.predict(x_test)
```

```
In [30]: #model evaluation  
decision_accuracy=round(accuracy_score(y_predict,y_test),3)  
decision_accuracy
```

```
Out[30]: 0.738
```

```
In [39]: #visualize decision tree  
plt.figure(figsize=(12,8))  
tree.plot_tree(decision_model,filled=True,feature_names=x.columns,class_names=['R','M'],label='all')  
plt.show()
```



```
In [40]: #random forest  
from sklearn.ensemble import RandomForestClassifier
```

```
In [42]: random_model=RandomForestClassifier(n_estimators=100,random_state=2)
```

```
In [43]: random_model.fit(x_train,y_train)
```

```
Out[43]: ▼ RandomForestClassifier  
RandomForestClassifier(random_state=2)
```

```
In [44]: y_predict=random_model.predict(x_test)
```

```
In [46]: #model evaluation  
random_accuracy=round(accuracy_score(y_predict,y_test),3)  
random_accuracy
```

```
Out[46]: 0.952
```

```
In [50]: #feature importance  
importance=random_model.feature_importances_  
feature_name=x.columns  
feature_imp_df=pd.DataFrame({'feature':feature_name,  
                             'importance':importance})  
print(feature_imp_df.sort_values(by='importance',ascending=False))
```

	feature	importance
10	10	0.062564
11	11	0.050623
8	8	0.045347
48	48	0.040870
47	47	0.037796
36	36	0.031665
50	50	0.027847
20	20	0.026632
46	46	0.024030
35	35	0.023441
22	22	0.023441
9	9	0.022913
45	45	0.019324
44	44	0.019167
34	34	0.018222
26	26	0.017482
30	30	0.017344
19	19	0.016905
38	38	0.016594
42	42	0.016578
27	27	0.016358
51	51	0.015585
18	18	0.014425
21	21	0.014414
12	12	0.014297
17	17	0.014212
39	39	0.014094
43	43	0.013720
14	14	0.013641
0	0	0.013421
23	23	0.013275
49	49	0.013040
16	16	0.013012
7	7	0.012779
54	54	0.011897
15	15	0.011760
31	31	0.011423
4	4	0.011127
33	33	0.011031
3	3	0.010441
5	5	0.010428
41	41	0.010375
6	6	0.010343
28	28	0.010343
25	25	0.010049
56	56	0.009748
24	24	0.009724
1	1	0.009583
59	59	0.009431
58	58	0.009372
57	57	0.009342
40	40	0.009191
52	52	0.009127
29	29	0.008620
55	55	0.007865
2	2	0.007401
13	13	0.007109
53	53	0.006803
37	37	0.006205
32	32	0.006203

```
In [51]: #naive bayes
from sklearn.naive_bayes import BernoulliNB
```

```
In [52]: naive_model=BernoulliNB()
```

```
In [53]: naive_model.fit(x_train,y_train)
```

```
Out[53]: ▼ BernoulliNB
BernoulliNB()
```

```
In [54]: y_predict=naive_model.predict(x_test)
```

```
In [55]: #model evaluation
naive_accuracy=round(accuracy_score(y_predict,y_test),3)
naive_accuracy
```

```
Out[55]: 0.524
```

```

In [57]: #k-nearest neighbour
from sklearn.neighbors import KNeighborsClassifier

In [58]: k=7
k_model=KNeighborsClassifier(n_neighbors=k)

In [60]: k_model.fit(x_train,y_train)

Out[60]:
▼      KNeighborsClassifier
KNeighborsClassifier(n_neighbors=7)

In [61]: y_predict=k_model.predict(x_test)

In [62]: #model evaluation
k_accuracy=round(accuracy_score(y_predict,y_test),3)
k_accuracy

Out[62]: 0.738

In [63]: model_evaluation=pd.DataFrame({'accuracy score':[logistic_accuracy,decision_accuracy,
                                                         random_accuracy,naive_accuracy,
                                                         k_accuracy]},index=['logistic reg','decision tree','random forest',
                                                         'naive bayes','knn'])

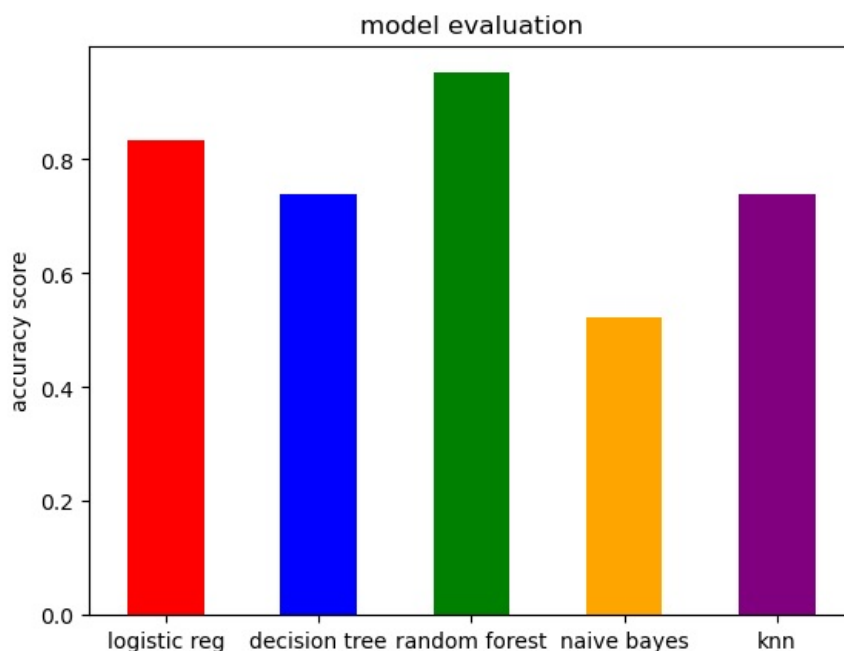
In [64]: model_evaluation

Out[64]:
      accuracy score
logistic reg      0.833
decision tree      0.738
random forest      0.952
naive bayes        0.524
knn                0.738

In [94]: colors=['red','blue','green','orange','purple']
model_evaluation.plot(kind='bar',y='accuracy score',legend=False,color=colors)
plt.ylabel('accuracy score')
plt.title('model evaluation')
plt.xticks(rotation=360)

Out[94]: (array([0, 1, 2, 3, 4]),
 [Text(0, 0, 'logistic reg'),
  Text(1, 0, 'decision tree'),
  Text(2, 0, 'random forest'),
  Text(3, 0, 'naive bayes'),
  Text(4, 0, 'knn')])

```



observation:- random forest can predict more accurately

as from the above chart it is clear that random forest is more accurate for predicting the roc/mine underwater, make a predictive system

```
In [97]: input_data=(0.0162,0.0041,0.0239,0.0441,0.0630,0.0921,0.1368,0.1078,0.1552,0.1779,0.2164,
0.2568,0.3089,0.3829,0.4393,0.5335,0.5996,0.6728,0.7309,0.8092,0.8941,0.9668,
1.0000,0.9893,0.9376,0.8991,0.9184,0.9128,0.7811,0.6018,0.3765,0.3300,0.2280,
0.0212,0.1117,0.1788,0.2373,0.2843,0.2241,0.2715,0.3363,0.2546,0.1867,0.2160,
0.1278,0.0768,0.1070,0.0946,0.0636,0.0227,0.0128,0.0173,0.0135,0.0114,0.0062,0.0157,0.0088,0.0036,0
#coping some values and using them to see if prediction is done correctly
#this should give-> M

#changing input data to numpy array for easy processing
data_array=np.asarray(input_data)

#reshape the data, as we are predicting for one instance and so model doesn't get confused with no. of data poi
data=data_array.reshape(1,-1)

#making prediction
prediction=random_model.predict(data)
prediction

if(prediction[0]=='R'):
    print('object is a rock')
else:
    print('object is a mine')
```

object is a mine

prediction is done correctly